

LEMANS SCHOOL OF AI | SESSION 4.04



SESSION ANIMÉE PAR
JEAN-BAPTISTE & PAUL

LEMANS
SCHOOL OF AI

ARTIFICIAL INTELLIGENCE STUDY GROUP MADE BY DOERS FOR DOERS IN LE MANS

SESSION D'APPROFONDISSEMENT

GENERATIVE ADVERSARIAL NETWORKS

JEUDI 05 NOVEMBRE 2020 | 18H30 | LE MANS INNOVATION

Sommaire

- Rappel principes des GAN
- Entraînement d'un GAN
- Evaluation d'un GAN
- StyleGAN

rappel

de quoi s'agit-il ?

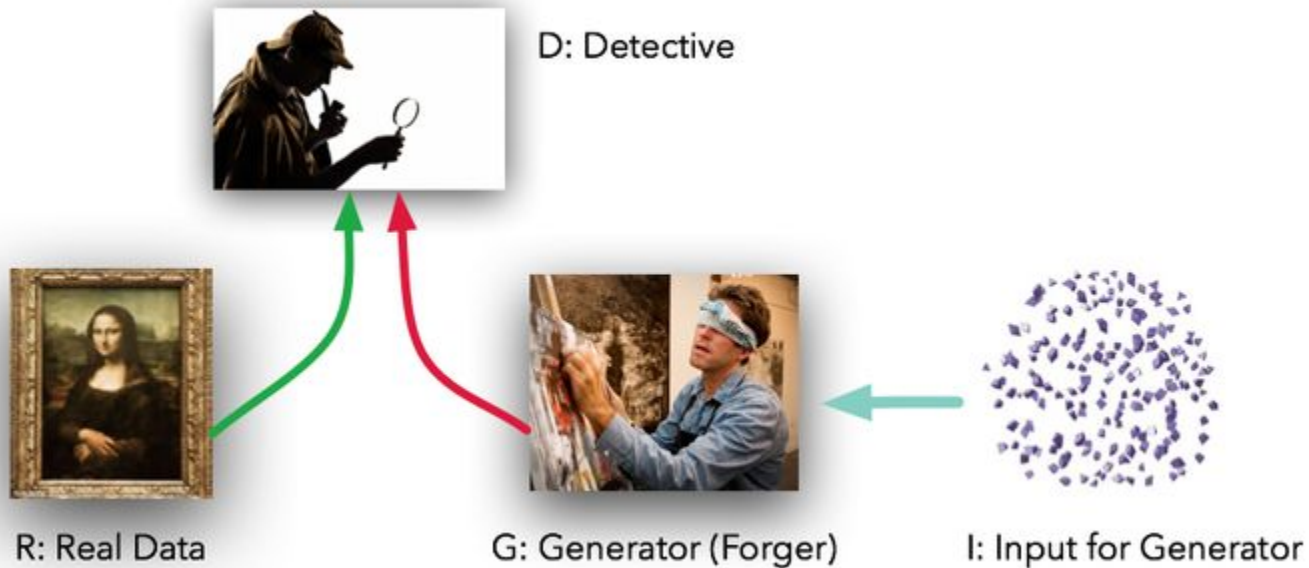
- **GAN** = type de modèle de DL

 la production de données fausses qui semblent réelles

- la particularité : entraînement non supervisé entre deux réseaux de neurones :

un **générateur** contre un **discriminateur**.

architecture d'un GAN



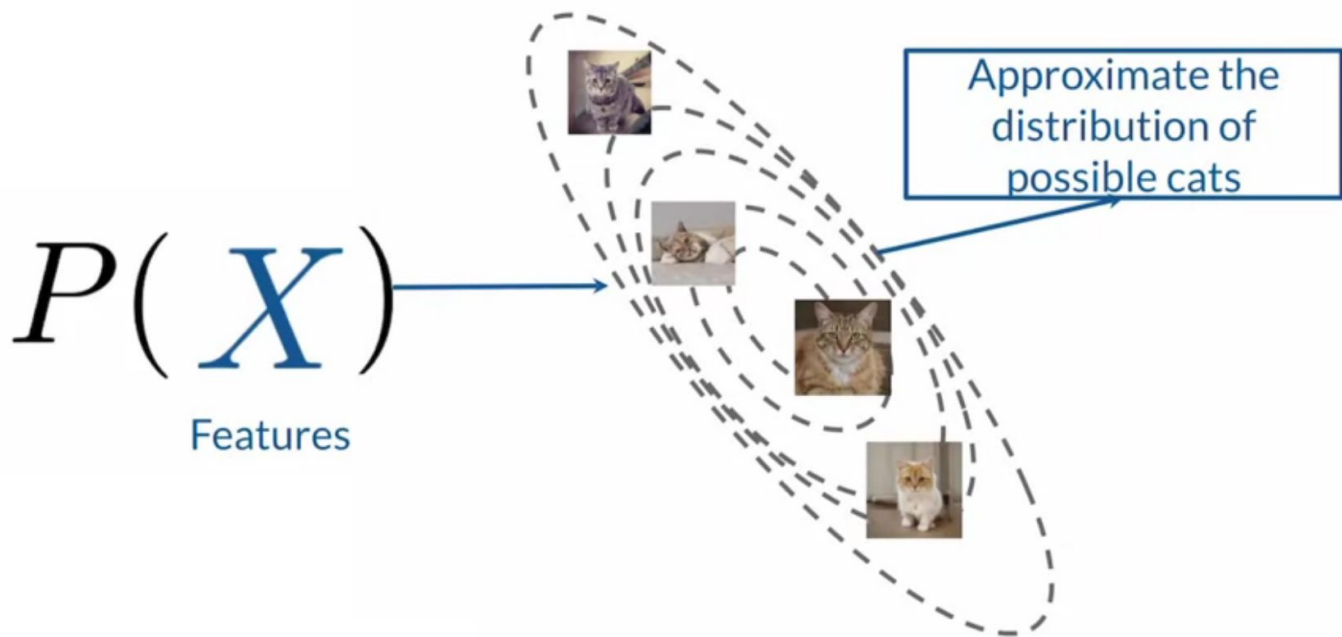
Discriminateur

un réseau de neurones de type classifieur


$$P(\underset{\text{Class}}{Y} \mid \underset{\text{Features}}{X})$$

$$P(\underset{\text{Class}}{\text{Fake}} \mid \underset{\text{Features}}{\img alt="A small, abstract, colorful image representing features, possibly a face or object, rendered in a heatmap style with blue, orange, and green." data-bbox="312 641 444 876}}) = 0.85 \longrightarrow \boxed{\text{Fake}}$$

Générateur



Training Data



Discriminator

1 (Real)
0 (Fake)

Latent Sample

-0.19972104, 0.42638235, -0.71335986,
0.27617624, 0.0455994, -0.82961057,
0.7449326, 0.305852, -0.81311934,
0.02522898, 0.60752668, 0.42092858,
0.06428853, 0.14700781, 0.42457545,
0.04710504, 0.26471074, -0.39863341,
-0.41719925, -0.71651508, 0.26192929,
-0.88549566, 0.65559716, -0.18518651,



Generator



Generated Image



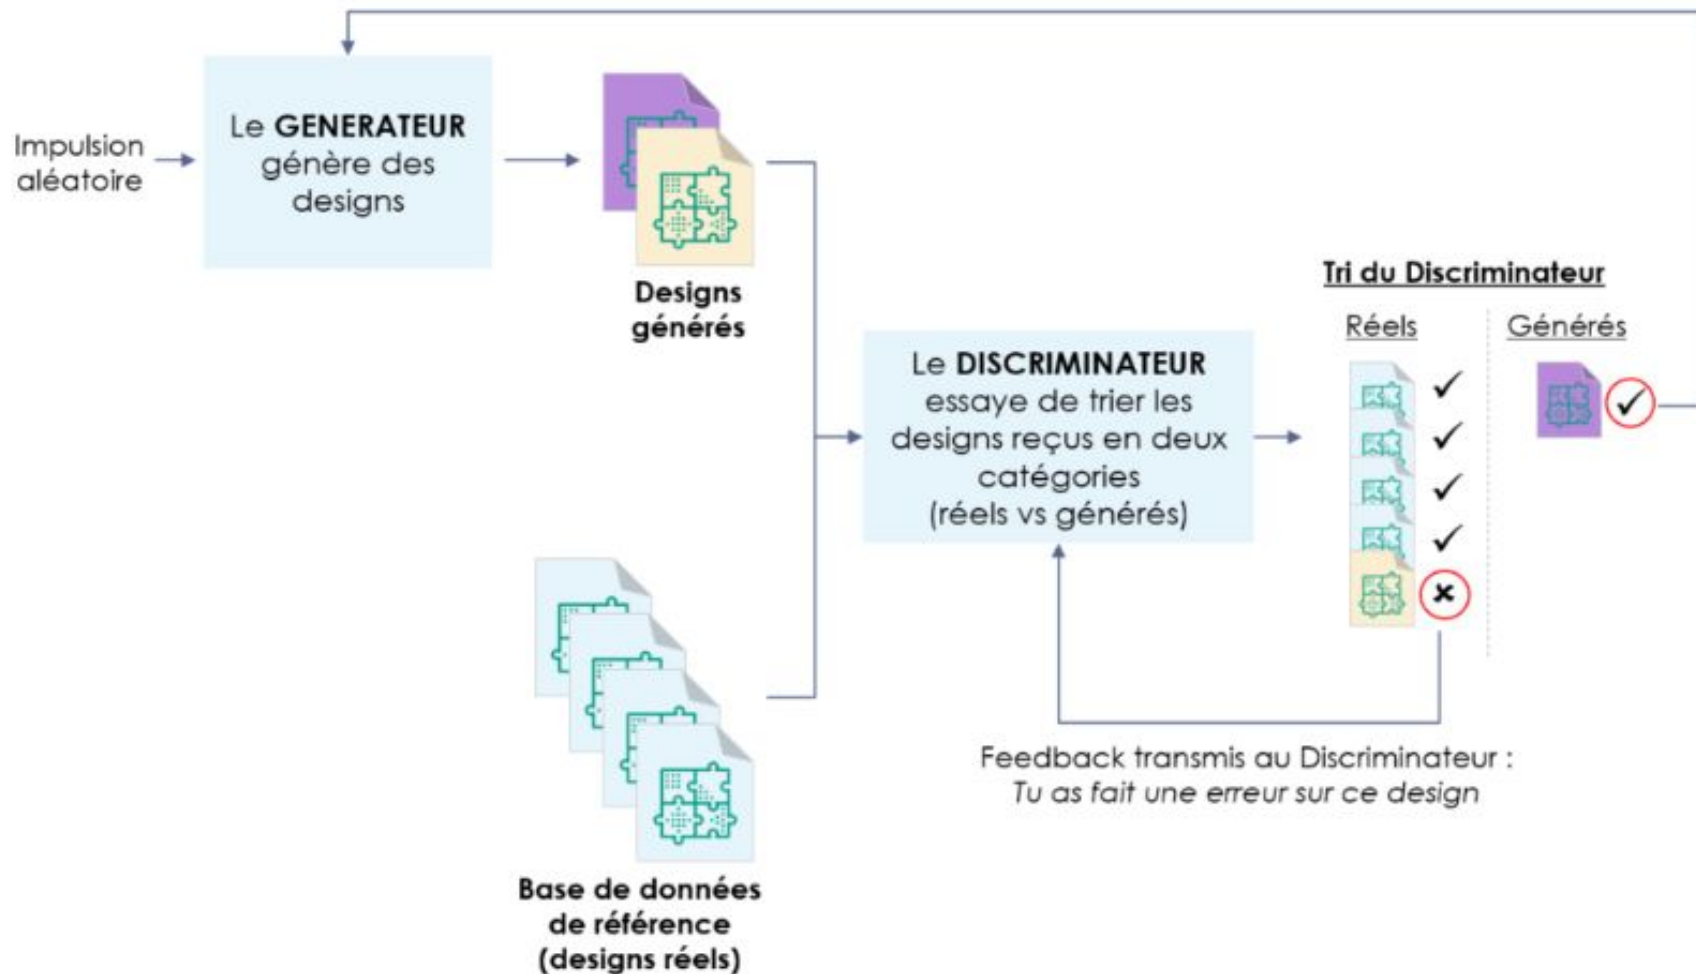
entrainement

Étapes d'un entraînement d'un GAN

- 1/ pré-entraînement du discriminateur
- 2/ entraînement du discriminateur
- 3/ entraînement du générateur

Feedback transmis au Générateur :


Tu n'as pas su tromper le Discriminateur sur tel design



2/ entraînement du discriminateur

- on fournit au réseau **discriminateur** les données contrefaites générées par le réseau générateur et il doit déterminer lesquelles sont vraies et lesquelles ne le sont pas.

Entraînement du discriminateur

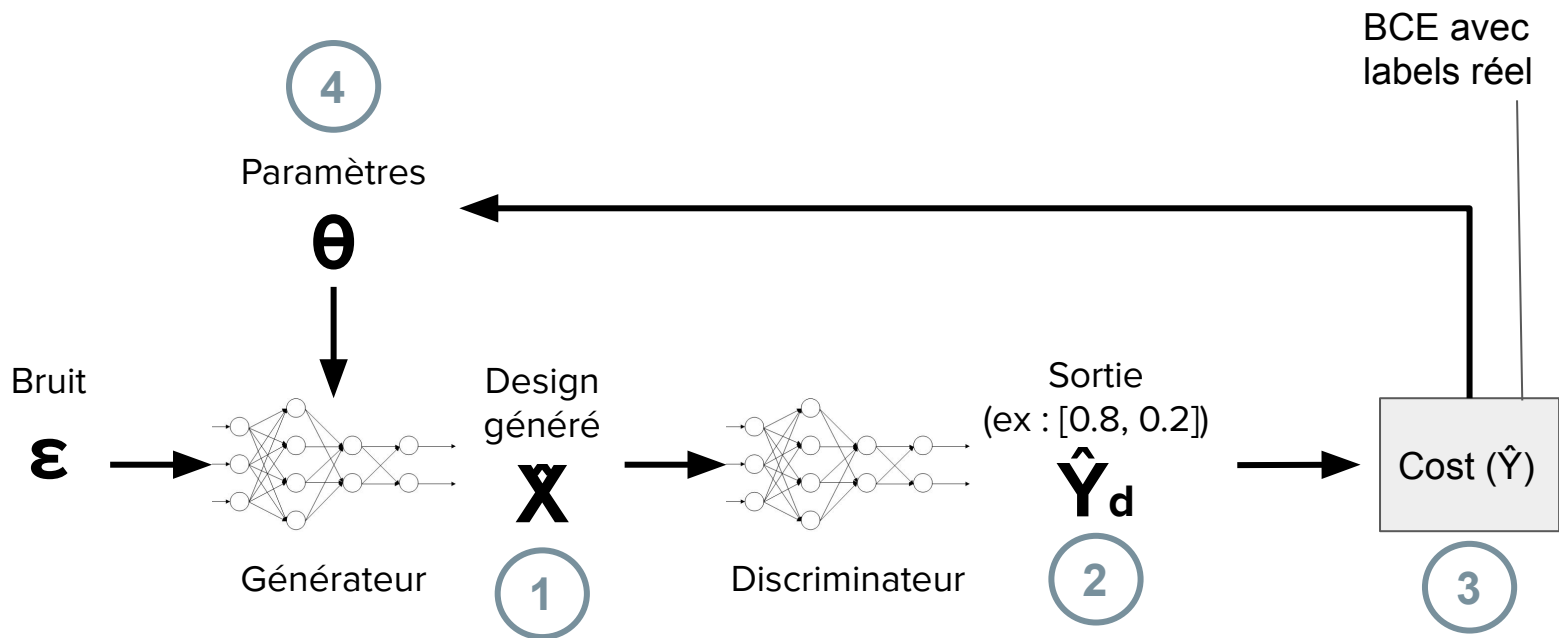


3/ entraînement du générateur

le réseau **générateur** s'améliore avec les résultats issus du deuxième entraînement du réseau discriminateur.

Le réseau **générateur** apprend à reconnaître les failles du discriminateur et cherche à les utiliser pour générer des ensembles de données contrefaites encore plus réalistes.

Entraînement du générateur



Objectif pour :

- le Générateur : $\hat{\mathbf{Y}}_d = [1, 0]$
- le Discriminateur : $\hat{\mathbf{Y}}_d = [0, 1]$

challenge

- le challenge :
 - les 2 réseaux doivent apprendre progressivement ensemble
 - trouver un équilibre
- si Discriminateur trop fort (prédiction 100% faux)


le Générateur n'apprend pas (pas de mise à jour des poids)
- si Générateur trop fort (prédiction 100% réel)

L'apprentissage est terminé


DCGAN

- pour améliorer les performances d'un modèle de génération d'image, le **Deep Convolutional GAN** est souvent utilisé
- Le **discriminateur** : un réseau à convolution
- le **générateur** : un réseau à convolution transposée ou “déconvolution”

Réseaux à convolution (discriminateur)




Réseaux à déconvolution (générateur)



DCGAN - exemple d'implémentation

- avec Pytorch
- dataset de chiffres manuscrits de MNIST
- objectif : générer une image synthétique d'un chiffre



DCGAN - le générateur

```
class Generator(nn.Module):
    """
    Generator Class
    Values:
        z_dim: the dimension of the noise vector, a scalar
        im_chan: the number of channels in the images, fitted for the dataset used, a scalar
                  (MNIST is black-and-white, so 1 channel is your default)
        hidden_dim: the inner dimension, a scalar
    """
    def __init__(self, z_dim=10, im_chan=1, hidden_dim=64):
        super(Generator, self).__init__()
        self.z_dim = z_dim
        # Build the neural network
        self.gen = nn.Sequential(
            self.make_gen_block(z_dim, hidden_dim * 4),
            self.make_gen_block(hidden_dim * 4, hidden_dim * 2, kernel_size=4, stride=1),
            self.make_gen_block(hidden_dim * 2, hidden_dim),
            self.make_gen_block(hidden_dim, im_chan, kernel_size=4, final_layer=True),
        )

    def make_gen_block(self, input_channels, output_channels, kernel_size=3, stride=2, final_layer=False):
        """
        Function to return a sequence of operations corresponding to a generator block of DCGAN,
        corresponding to a transposed convolution, a batchnorm (except for in the last layer), and an activation.
        Parameters:
            input_channels: how many channels the input feature representation has
            output_channels: how many channels the output feature representation should have
            kernel_size: the size of each convolutional filter, equivalent to (kernel_size, kernel_size)
            stride: the stride of the convolution
            final_layer: a boolean, true if it is the final layer and false otherwise
                        (affects activation and batchnorm)
        """
        # Build the neural block
        if not final_layer:
            return nn.Sequential(
                nn.ConvTranspose2d(input_channels, output_channels, kernel_size, stride=stride),
                nn.BatchNorm2d(output_channels),
                nn.ReLU(inplace=True)
            )
```

DCGAN

- le

discrim

inateur

```
class Discriminator(nn.Module):
    """
    Discriminator Class
    Values:
        im_chan: the number of channels in the images, fitted for the dataset used, a scalar
                  (MNIST is black-and-white, so 1 channel is your default)
        hidden_dim: the inner dimension, a scalar
    """
    def __init__(self, im_chan=1, hidden_dim=16):
        super(Discriminator, self).__init__()
        self.disc = nn.Sequential(
            self.make_disc_block(im_chan, hidden_dim),
            self.make_disc_block(hidden_dim, hidden_dim * 2),
            self.make_disc_block(hidden_dim * 2, 1, final_layer=True),
        )

    def make_disc_block(self, input_channels, output_channels, kernel_size=4, stride=2, final_layer=False):
        """
        Function to return a sequence of operations corresponding to a discriminator block of DCGAN,
        corresponding to a convolution, a batchnorm (except for in the last layer), and an activation.
        Parameters:
            input_channels: how many channels the input feature representation has
            output_channels: how many channels the output feature representation should have
            kernel_size: the size of each convolutional filter, equivalent to (kernel_size, kernel_size)
            stride: the stride of the convolution
            final_layer: a boolean, true if it is the final layer and false otherwise
                        (affects activation and batchnorm)
        """

        # Build the neural block
        if not final_layer:
            return nn.Sequential(
                nn.Conv2d(input_channels, output_channels, kernel_size, stride=stride),
                nn.BatchNorm2d(output_channels),
                nn.LeakyReLU(0.2)
            )
        else: # Final Layer
            return nn.Sequential(
                nn.Conv2d(input_channels, output_channels, kernel_size, stride=stride)
```

DCGAN - parametre entraînement

```
criterion = nn.BCEWithLogitsLoss()
z_dim = 64
display_step = 500
batch_size = 128
# A learning rate of 0.0002 works well on DCGAN
lr = 0.0002

# These parameters control the optimizer's momentum, which you can read more about here:
# https://distill.pub/2017/momentum/
beta_1 = 0.5
beta_2 = 0.999
device = 'cuda'

# You can tranform the image values to be between -1 and 1 (the range of the tanh activation)
transform = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.5,), (0.5,)),
])

dataloader = DataLoader(
    MNIST('.', download=False, transform=transform),
    batch_size=batch_size,
    shuffle=True)
```


DCCGAN - initialisation

```
gen = Generator(z_dim).to(device)
gen_opt = torch.optim.Adam(gen.parameters(), lr=lr, betas=(beta_1, beta_2))
disc = Discriminator().to(device)
disc_opt = torch.optim.Adam(disc.parameters(), lr=lr, betas=(beta_1, beta_2))

# You initialize the weights to the normal distribution
# with mean 0 and standard deviation 0.02
def weights_init(m):
    if isinstance(m, nn.Conv2d) or isinstance(m, nn.ConvTranspose2d):
        torch.nn.init.normal_(m.weight, 0.0, 0.02)
    if isinstance(m, nn.BatchNorm2d):
        torch.nn.init.normal_(m.weight, 0.0, 0.02)
        torch.nn.init.constant_(m.bias, 0)
gen = gen.apply(weights_init)
disc = disc.apply(weights_init)
```


DCGAN - entraine ment

```
n_epochs = 50
cur_step = 0
mean_generator_loss = 0
mean_discriminator_loss = 0
for epoch in range(n_epochs):
    # Dataloader returns the batches
    for real, _ in tqdm(dataloader):
        cur_batch_size = len(real)
        real = real.to(device)

        ## Update discriminator ##
        disc_opt.zero_grad()
        fake_noise = get_noise(cur_batch_size, z_dim, device=device)
        fake = gen(fake_noise)
        disc_fake_pred = disc(fake.detach())
        disc_fake_loss = criterion(disc_fake_pred, torch.zeros_like(disc_fake_pred))
        disc_real_pred = disc(real)
        disc_real_loss = criterion(disc_real_pred, torch.ones_like(disc_real_pred))
        disc_loss = (disc_fake_loss + disc_real_loss) / 2

        # Keep track of the average discriminator loss
        mean_discriminator_loss += disc_loss.item() / display_step
    # Update gradients
    disc_loss.backward(retain_graph=True)
    # Update optimizer
    disc_opt.step()
```

DCGAN - entaine ment (suite)

```
# Keep track of the average discriminator loss
mean_discriminator_loss += disc_loss.item() / display_step

# Update gradients
disc_loss.backward(retain_graph=True)

# Update optimizer
disc_opt.step()


## Update generator ##
gen_opt.zero_grad()
fake_noise_2 = get_noise(cur_batch_size, z_dim, device=device)
fake_2 = gen(fake_noise_2)
disc_fake_pred = disc(fake_2)
gen_loss = criterion(disc_fake_pred, torch.ones_like(disc_fake_pred))
gen_loss.backward()
gen_opt.step()

# Keep track of the average generator loss
mean_generator_loss += gen_loss.item() / display_step


## Visualization code ##
if cur_step % display_step == 0 and cur_step > 0:
    print(f"Step {cur_step}: Generator loss: {mean_generator_loss}, discriminator loss: {mean_discriminator_loss}")
    show_tensor_images(fake)
    show_tensor_images(real)
    mean_generator_loss = 0
    mean_discriminator_loss = 0
cur_step += 1
```

DCGAN - résultat (suite)

Step 1000: Generator loss: 2.0383475271463407, discriminator loss: 0.24150314955413335



Step 23000: Generator loss: 0.6993783376812931, discriminator loss: 0.6958930463790896





Remarques pour l'amélioration du modèle

- Utilisation de \tanh pour la couche output du générateur (Scale the image pixel value between -1 and 1.)
- Batch Normalization aide à la stabilisation de l'entraînement
- Eviter le max pooling
- utiliser la fonction d'optimisation Adam


Problème de l'effet damier

- Problème courant pour les **déconvolutions**
- l'utilisation d'un upsampling suivi d'une convolution est une technique de plus en plus populaire aujourd'hui pour éviter ce problème de damier.

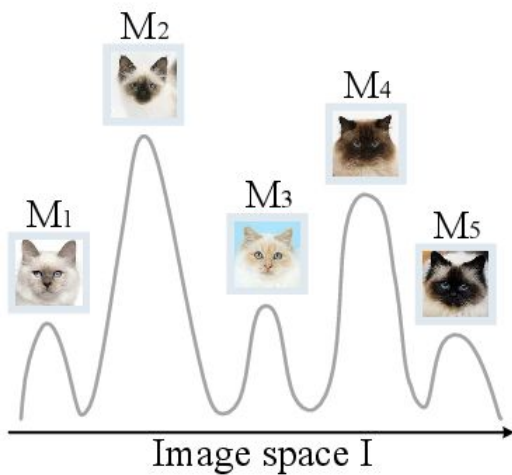


Mode collapse

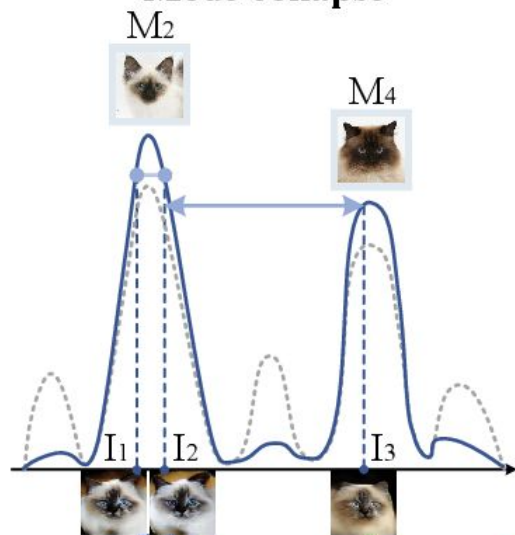
- problème fréquemment rencontré
- observable ici par une représentation d'un sous-ensemble des chiffres (ici 9, 7, 1)




Real data



Mode collapse




Mode seeking



$$\frac{d_I(I_a, I_b)}{d_Z(z_a, z_b)}$$

Legend for the ratio values:

- 0.68 (green double-headed arrow)
- 0.58 (green line with dots)
- 0.62 (blue double-headed arrow)
- 0.17** (blue line with dots)




Latent space Z

Problème Vanishing gradient

- c'est un problème lié à l'utilisation la fonction de coût BCE (Binary Cross Entropy)
- c'est un problème qui survient quand le **Discriminateur** s'améliore trop vite



Problème Vanishing gradient (suite)

objectif du générateur : tendre vers la distribution réelle





Problème Vanishing gradient (suite)

gradient non-zero = le générateur s'améliore




Problème Vanishing gradient (suite)

gradient proche de zéro = le générateur ne peut plus s'améliorer



W-loss

- fonction de coût alternative : Wasserstein loss
- approximatation de la Earth mover's distance



W-loss

- résout les problèmes :
 - mode collapse
 - vanished gradient
- implique souvent la mise en place de pénalités sur les poids lors de la rétropropagation

Exemples de programme illustrant des modèles instables :

<https://www.aiproblog.com/index.php/2019/07/07/how-to-identify-and-diagnose-gan-failure-modes/>

Évaluation

Évaluation

- Inception-v3 and Embeddings
- Fréchet Inception Distance (FID)

Inception-v3 architecture


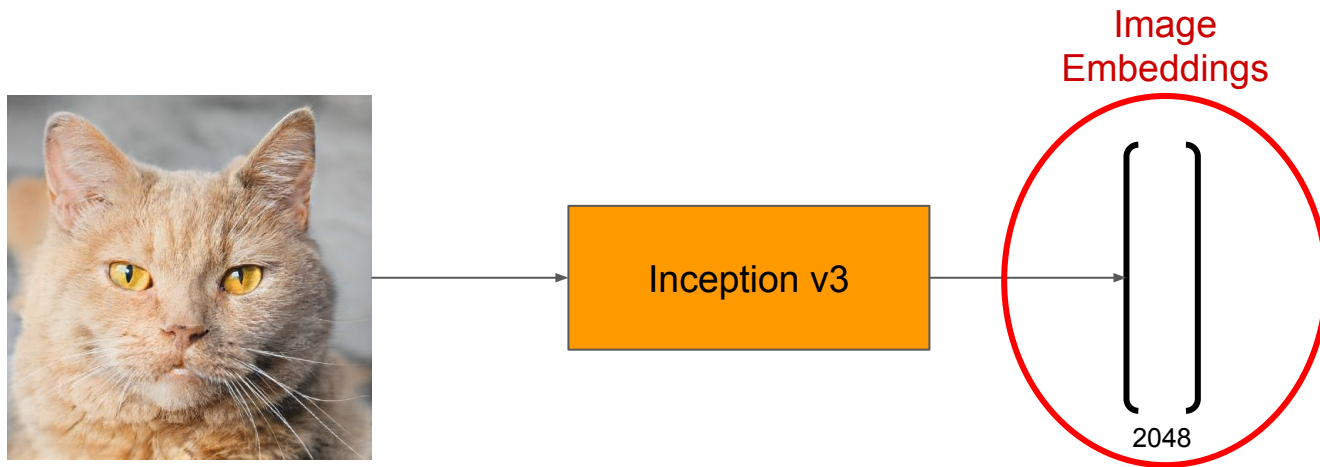




Image Embeddings



Fréchet distance



Fréchet distance 2 dimensions



$$d(X, Y) =$$

$$(\mu_X - \mu_Y)^2 + (\sigma_X - \sigma_Y)^2$$

Moyenne

Écart-type

Fréchet distance 2 dimension to multidimension

$$(\mu_X - \mu_Y)^2 + (\sigma_X - \sigma_Y)^2$$



$$\|\mu_X - \mu_Y\|^2 + \text{Tr} \left(\Sigma_X + \Sigma_Y - 2\sqrt{\Sigma_X \Sigma_Y} \right)$$

Fréchet inception distance (FID)

$$\text{FID} = \|\mu_X - \mu_Y\|^2 + \text{Tr} \left(\Sigma_X + \Sigma_Y - 2\sqrt{\Sigma_X \Sigma_Y} \right)$$

X : Reel image Embeddings

(ex : 50 000 vecteurs de dimension 2048)

Y : Fake image Embeddings

(ex : 50 000 vecteurs de dimension 2048)

Conclusion :

Plus FID est petit plus les images sont proches donc plus FID est petit meilleur est notre générateur.

FID limits

- Utilisation d'un modèle inception pré-entraîné qui ne va pas forcément récupérer toutes les features des images. (dépend de la proximité de son dataset d'entraînement avec les images à évaluer)
- A besoin d'un large échantillon d'image pour éviter un maximum le bruit. (Au moins 50 000 images)
- Lent à exécuter

Pour résumer

- FID calcule la distance entre les image réelles et fakes
- FID utilise inception et la distance de Fréchet
- FID a besoin d'un large échantillon pour bien fonctionner

StyleGAN

Et Aujourd'hui, quel est l'état de l'art?



2014



2015



2016




2017



2018


StyleGAN




StyleGAN

- Objectifs de StyleGAN
 - Plus grande fidélité
 - Plus grande diversité
 - Plus de contrôle sur les features
- Composants de StyleGAN :
 - Progressive Growing
 - Adaptive Instance Normalization (AdaIN)
 - Noise mapping network
 - Style mixing


Progressive growing




...




Progressive growing : Generator



Progressive growing : Discriminator



Progressive growing in context



Progressive growing Implementation

```
class MicroStyleGANGeneratorBlock(nn.Module):
```

```
...
```

Micro StyleGAN Generator Block Class

Values:

in_chan: the number of channels in the input, a scalar
out_chan: the number of channels wanted in the output, a scalar
w_dim: the dimension of the intermediate noise vector, a scalar
kernel_size: the size of the convolving kernel
starting_size: the size of the starting image

```
...
```

```
def __init__(self, in_chan, out_chan, w_dim, kernel_size, starting_size, use_upsample=True):
```

```
    super().__init__()
```

```
    self.use_upsample = use_upsample
```

```
    # 1. Upsample to the starting size, bilinearly (https://pytorch.org/docs/master/generated/torch.nn.Upsample.html)
```

```
    # 2. Create a kernel_size convolution which takes in
```

```
    # an image with in_chan and outputs one with out_chan (https://pytorch.org/docs/stable/generated/torch.nn.Conv2d.html)
```

```
    # 3. Create an object to inject noise
```

```
    # 4. Create an AdaIN object
```

```
    # 5. Create a LeakyReLU activation with slope 0.2
```

```
    if self.use_upsample:
```

```
        self.upsample = nn.Upsample((starting_size, starting_size), mode='bilinear')
```

```
    self.conv = nn.Conv2d(in_chan, out_chan, kernel_size, padding=1) # Padding is used to maintain the image size
```

```
    self.inject_noise = InjectNoise(out_chan)
```

```
    self.adain = AdaIN(out_chan, w_dim)
```

```
    self.activation = nn.LeakyReLU(0.2)
```

```
def forward(self, x, w):
```

```
...
```

Function for completing a forward pass of MicroStyleGANGeneratorBlock: Given an x and w,

computes a StyleGAN generator block.

Parameters:

x: the input into the generator, feature map of shape (n_samples, channels, width, height)

w: the intermediate noise vector

```
...
```

```
if self.use_upsample:
```

```
    x = self.upsample(x)
```

```
x = self.conv(x)
```

```
x = self.inject_noise(x)
```

```
x = self.activation(x)
```

```
x = self.adain(x, w)
```

```
return x
```

```
class MicroStyleGANGenerator(nn.Module):
```

```
...
```

Micro StyleGAN Generator Class

Values:

z_dim: the dimension of the noise vector, a scalar

map_hidden_dim: the mapping inner dimension, a scalar

w_dim: the dimension of the intermediate noise vector, a scalar

in_chan: the dimension of the constant input, usually w_dim, a scalar

out_chan: the number of channels wanted in the output, a scalar

kernel_size: the size of the convolving kernel

hidden_chan: the inner dimension, a scalar

```
...
```

```
def __init__(self, z_dim, map_hidden_dim, w_dim, in_chan, out_chan, kernel_size, hidden_chan):
```

```
    super().__init__()
```

```
    self.map = MappingLayers(z_dim, map_hidden_dim, w_dim)
```

```
    # Typically this constant is initiated to all ones, but you will initiate to a
```

```
    # Gaussian to better visualize the network's effect
```

```
    self.starting_constant = nn.Parameter(torch.randn(1, in_chan, 4, 4))
```

```
    self.block0 = MicroStyleGANGeneratorBlock(in_chan, hidden_chan, w_dim, kernel_size, 4, use_upsample=False)
```

```
    self.block1 = MicroStyleGANGeneratorBlock(hidden_chan, hidden_chan, w_dim, kernel_size, 8)
```

```
    self.block2 = MicroStyleGANGeneratorBlock(hidden_chan, hidden_chan, w_dim, kernel_size, 16)
```

```
    # You need to have a way of mapping from the output noise to an image,
```

```
    # so you learn a 1x1 convolution to transform the e.g. 512 channels into 3 channels
```

```
    # (Note that this is simplified, with clipping used in the real StyleGAN)
```

```
    self.block1_to_image = nn.Conv2d(hidden_chan, out_chan, kernel_size=1)
```

```
    self.block2_to_image = nn.Conv2d(hidden_chan, out_chan, kernel_size=1)
```

```
    self.alpha = 0.2
```

```
def upsample_to_match_size(self, smaller_image, bigger_image):
```

```
...
```

Function for upsampling an image to the size of another: Given a two images (smaller and bigger),

upsamples the first to have the same dimensions as the second.

Parameters:

smaller_image: the smaller image to upsample

bigger_image: the bigger image whose dimensions will be upsampled to

```
...
```

```
return F.interpolate(smaller_image, size=bigger_image.shape[-2:], mode='bilinear')
```

```
def forward(self, noise, return_intermediate=False):
```

```
...
```

Function for completing a forward pass of MicroStyleGANGenerator: Given noise,

computes a StyleGAN iteration.

Parameters:

noise: a noise tensor with dimensions (n_samples, z_dim)

return_intermediate: a boolean, true to return the images as well (for testing) and false otherwise

```
...
```

```
x = self.starting_constant
```

```
w = self.map(noise)
```

```
x = self.block0(x, w)
```

```
x_small = self.block1(x, w) # First generator run output
```

```
x_small_image = self.block1_to_image(x_small)
```

```
x_big = self.block2(x_small, w) # Second generator run output
```

```
x_big_image = self.block2_to_image(x_big)
```

```
# Upsample first generator run output to be same size as second generator run output
```

```
x_small_upsample = self.upsample_to_match_size(x_small_image, x_big_image)
```

```
# Interpolate between the upsampled image and the image from the generator using alpha
```

```
interpolation = torch.lerp(x_small_upsample, x_big_image, self.alpha)
```

```
if return_intermediate:
```


```
    return interpolation, x_small_upsample, x_big_image
```

```
return interpolation
```


En résumé

- Progressive growing double progressivement la taille des images générées
- Progressive growing entraîne un entraînement plus rapide, plus stable et une meilleur résolution d'image

Noise Mapping Network




Noise Mapping Network



More possible to control single output features

Noise Mapping Network + Progressive Growing



Noise Mapping Network implementation

```
class MappingLayers(nn.Module):
    """
    Mapping Layers Class
    Values:
        z_dim: the dimension of the noise vector, a scalar
        hidden_dim: the inner dimension, a scalar
        w_dim: the dimension of the intermediate noise vector, a scalar
    """


    def __init__(self, z_dim, hidden_dim, w_dim):
        super().__init__()
        self.mapping = nn.Sequential(
            nn.Linear(z_dim, hidden_dim),
            nn.ReLU(),
            nn.Linear(hidden_dim, hidden_dim),
            nn.ReLU(),
            nn.Linear(hidden_dim, w_dim)
        )

    def forward(self, noise):
        """
        Function for completing a forward pass of MappingLayers:
        Given an initial noise tensor, returns the intermediate noise tensor.
        Parameters:
            noise: a noise tensor with dimensions (n_samples, z_dim)
        """
        return self.mapping(noise)
```


En résumé

- Noise Mapping permet de mieux contrôler les features
- W le vecteur de bruit intermédiaire est injecté à différents niveau dans le générateur

Adaptive Instance Normalization (AdaIN)




Adaptive Instance Normalization (AdaIN)



Etape 1 : Normaliser la sortie des couches de convolution en utilisant l'**Instance Normalisation**

Adaptive Instance Normalization (AdaIN)



$$\text{AdaIN}(\mathbf{x}_i, \mathbf{y}) = \mathbf{y}_{s,i} \frac{\mathbf{x}_i - \mu(\mathbf{x}_i)}{\sigma(\mathbf{x}_i)} + \mathbf{y}_{b,i}$$

Étape 1 : Instance normalization

Etape 2 : Appliquer l'**Adaptive Style** en utilisant W le vecteur de bruit intermédiaire

AdaIN implementation

```
class AdaIN(nn.Module):
    """
    AdaIN Class
    Values:
        channels: the number of channels the image has, a scalar
        w_dim: the dimension of the intermediate noise vector, a scalar
    """

    def __init__(self, channels, w_dim):
        super().__init__()

        # Normalize the input per-dimension
        self.instance_norm = nn.InstanceNorm2d(channels)

        # You want to map w to a set of style weights per channel.
        # Replace the Nones with the correct dimensions - keep in mind that
        # both linear maps transform a w vector into style weights
        # corresponding to the number of image channels.
        self.style_scale_transform = nn.Linear(w_dim, channels)
        self.style_shift_transform = nn.Linear(w_dim, channels)

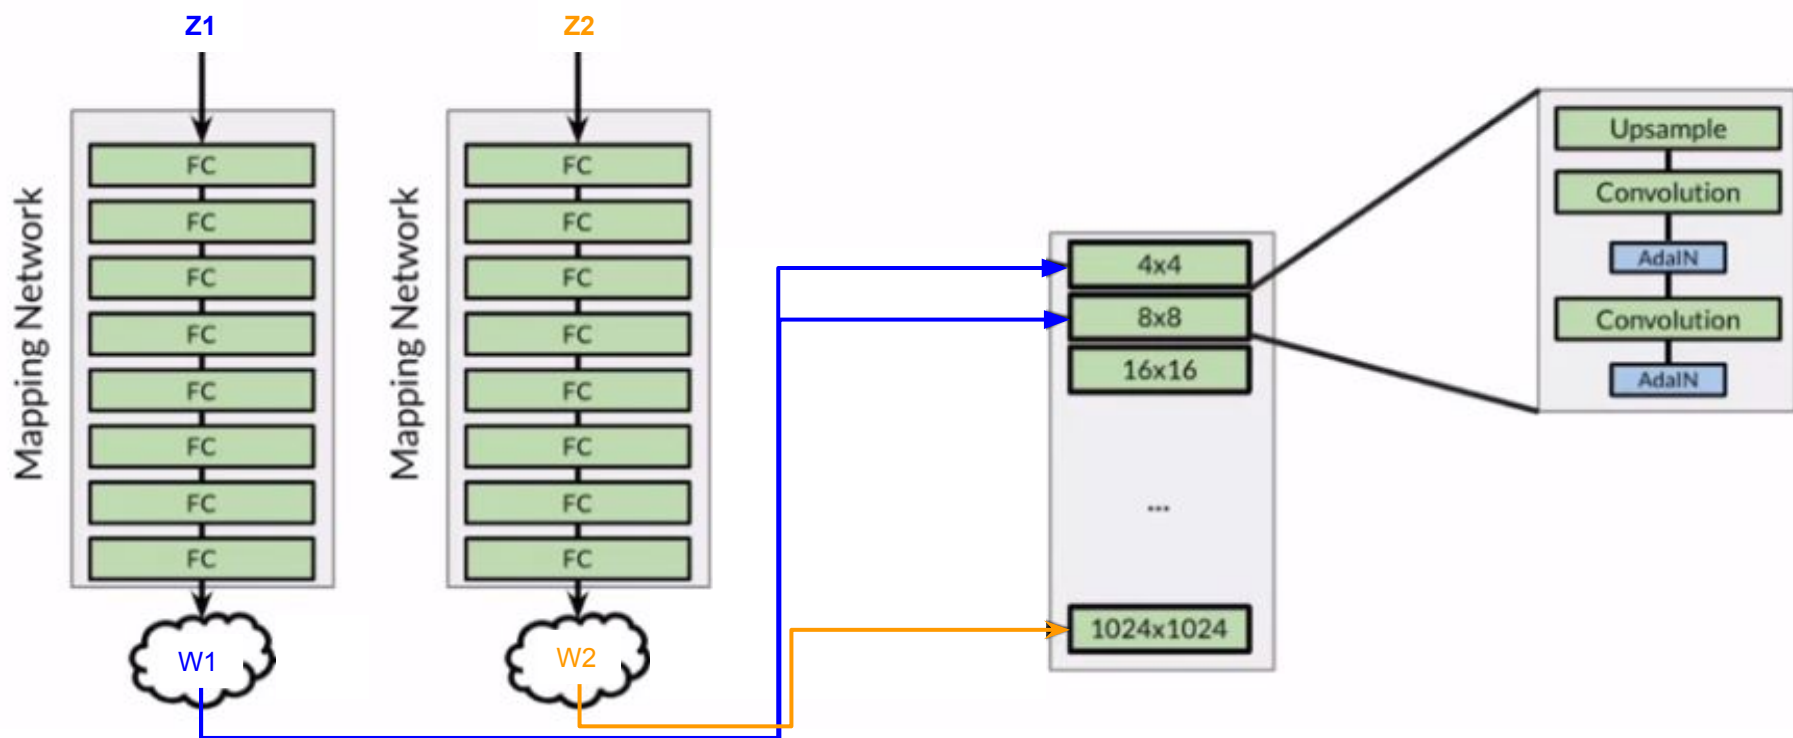
    def forward(self, image, w):
        """
        Function for completing a forward pass of AdaIN: Given an image and intermediate noise vector w,
        returns the normalized image that has been scaled and shifted by the style.
        Parameters:
            image: the feature map of shape (n_samples, channels, width, height)
            w: the intermediate noise vector
        """
        normalized_image = self.instance_norm(image)
        style_scale = self.style_scale_transform(w)[: , :, None, None]
        style_shift = self.style_shift_transform(w)[: , :, None, None]

        # Calculate the transformed image
        transformed_image = style_scale * normalized_image + style_shift
        return transformed_image
```


En résumé

- Instance Normalization est utilisé pour normaliser individuellement les exemples avant de leur appliquer un style
- AdaIN transfère les informations de style du vecteur de bruit intermédiaire w vers l'image générée.

Style Mixing



Style Mixing

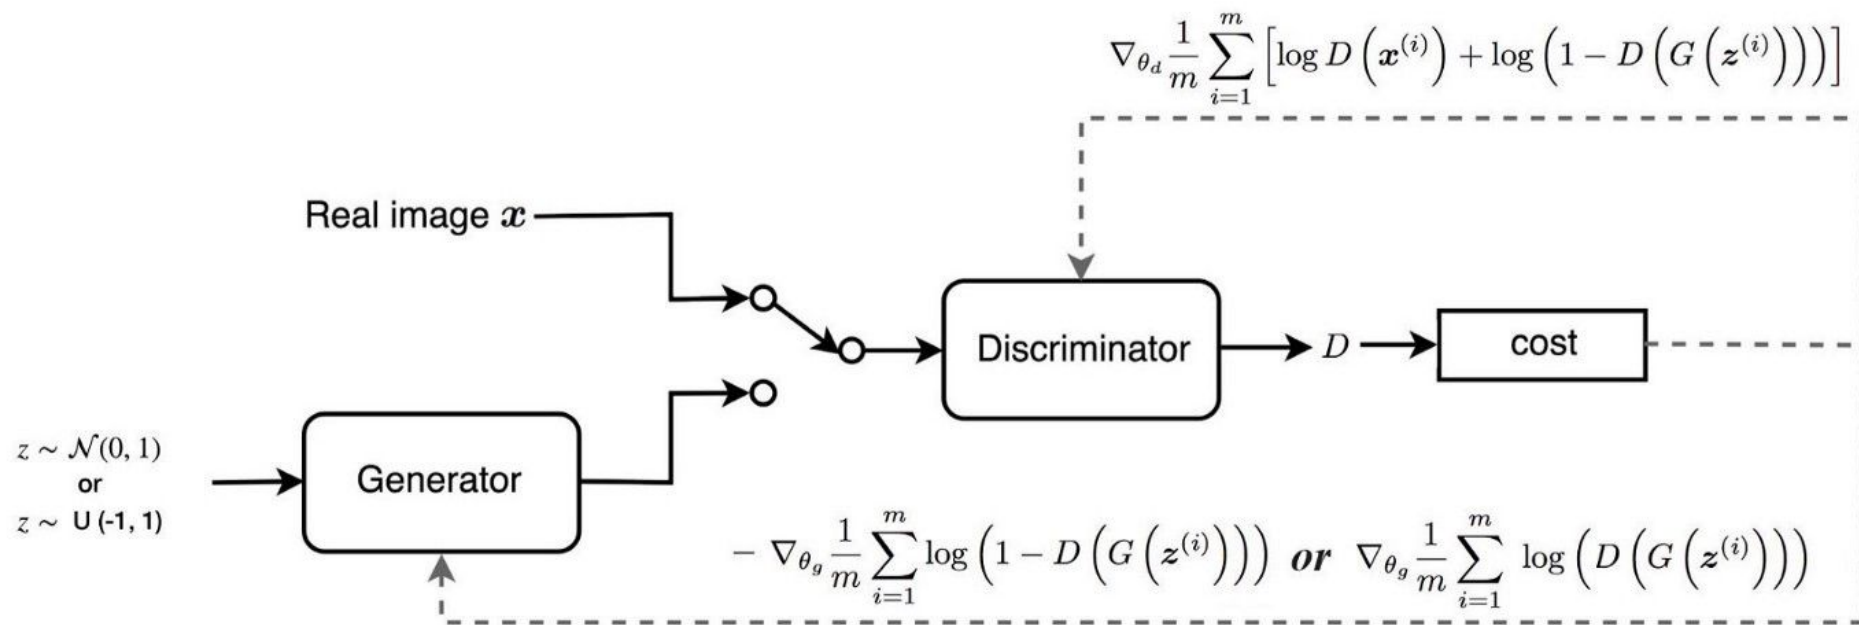


Merci !

Annexe

collection of implementation gan

<https://github.com/eriklindernoren/Keras-GAN>



Binary cross-entropy cost function

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m \left[\overset{\text{real } y^{(i)}=1}{\boxed{y^{(i)} \log h(x^{(i)}, \theta)}} + \overset{\text{fake } y^{(i)}=0}{\boxed{(1 - y^{(i)}) \log(1 - h(x^{(i)}, \theta))}} \right]$$

Prediction
Label
Features
Parameters

$h(x^{(i)}, \theta)$	$\log h(x^{(i)}, \theta)$	$\log(1 - h(x^{(i)}, \theta))$
~ 0	$-\infty$	~ 0
~ 1	~ 0	$-\infty$

Objectif pour :

- le Générateur : $J(\theta)$ plus grand possible
- le Discriminateur : $J(\theta)$ plus proche de 0 possible